

R-Laboratory 6

MANOJ KUMAR - 2048015

26/02/2021

1. Load mtcars dataset.

Loading mtcars dataset.

```
data(mtcars)
```

```
mtcars
```

##	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
## Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
## Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
## Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
## Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
## Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
## Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
## Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
## Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
## Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
## Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
## Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
## Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
## Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
## Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
## Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
## Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
## Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
## Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
## Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
## Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
## Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
## Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
## AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
## Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
## Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
## Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
## Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
## Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
## Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
## Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
## Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
## Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

2.install ridge and glmnet packages.

```
#install.packages(ridge)
#install.packages(glmnet)

library(ridge)      # Linear and Logistic ridge regression functions.
library(Matrix)
library(glmnet)     # Lasso and Elastic-Net Regularized Generalized Linear
Models

## Loaded glmnet 4.1
```

3.Perform the exploratory data analysis.

```
# Pre-processing EDA
df = mtcars

str(df)

## 'data.frame':    32 obs. of  11 variables:
## $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
## $ cyl : num   6  6  4  6  8  6  8  4  4  6 ...
## $ disp: num  160 160 108 258 360 ...
## $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
## $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
## $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
## $ qsec: num   16.5 17 18.6 19.4 17 ...
## $ vs  : num    0  0  1  1  0  1  0  1  1  1 ...
## $ am  : num    1  1  1  0  0  0  0  0  0  0 ...
## $ gear: num    4  4  4  3  3  3  3  4  4  4 ...
## $ carb: num    4  4  1  1  2  1  4  2  2  4 ...
```

Insight

- Totally, 32 observations of 11 variables.
- All the 11 features are numerical datatypes.

Summary

```
summary(df)
```

##	mpg	cyl	disp	hp
##	Min. :10.40	Min. :4.000	Min. : 71.1	Min. : 52.0
##	1st Qu.:15.43	1st Qu.:4.000	1st Qu.:120.8	1st Qu.: 96.5
##	Median :19.20	Median :6.000	Median :196.3	Median :123.0
##	Mean :20.09	Mean :6.188	Mean :230.7	Mean :146.7
##	3rd Qu.:22.80	3rd Qu.:8.000	3rd Qu.:326.0	3rd Qu.:180.0
##	Max. :33.90	Max. :8.000	Max. :472.0	Max. :335.0
##	drat	wt	qsec	vs
##	Min. :2.760	Min. :1.513	Min. :14.50	Min. :0.0000
##	1st Qu.:3.080	1st Qu.:2.581	1st Qu.:16.89	1st Qu.:0.0000
##	Median :3.695	Median :3.325	Median :17.71	Median :0.0000
##	Mean :3.597	Mean :3.217	Mean :17.85	Mean :0.4375

```
## 3rd Qu.:3.920 3rd Qu.:3.610 3rd Qu.:18.90 3rd Qu.:1.0000
## Max. :4.930 Max. :5.424 Max. :22.90 Max. :1.0000
##      am      gear      carb
## Min. :0.0000 Min. :3.000 Min. :1.000
## 1st Qu.:0.0000 1st Qu.:3.000 1st Qu.:2.000
## Median :0.0000 Median :4.000 Median :2.000
## Mean :0.4062 Mean :3.688 Mean :2.812
## 3rd Qu.:1.0000 3rd Qu.:4.000 3rd Qu.:4.000
## Max. :1.0000 Max. :5.000 Max. :8.000
```

Checking for missing values.

```
colSums(is.na(df))
```

```
## mpg cyl disp hp drat wt qsec vs am gear carb
## 0 0 0 0 0 0 0 0 0 0 0 0
```

#Checking for Empty Values

```
colSums(df=='')
```

```
## mpg cyl disp hp drat wt qsec vs am gear carb
## 0 0 0 0 0 0 0 0 0 0 0 0
```

#Checking for Duplicate values

```
library(tidyverse)
```

```
## — Attaching packages
```

```
tidyverse 1.3.0 —
```

```
## ✓ ggplot2 3.3.2 ✓ purrr 0.3.4
## ✓ tibble 3.0.3 ✓ dplyr 1.0.4
## ✓ tidyr 1.1.1 ✓ stringr 1.4.0
## ✓ readr 1.4.0 ✓ forcats 0.5.1
```

```
## — Conflicts
```

```
tidyverse_conflicts() —
```

```
## x tidyr::expand() masks Matrix::expand()
## x dplyr::filter() masks stats::filter()
## x dplyr::lag() masks stats::lag()
## x tidyr::pack() masks Matrix::pack()
## x tidyr::unpack() masks Matrix::unpack()
```

```
duplicated(df)
```

```
## [1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
## [13] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

```
FALSE
## [25] FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

Insight

- The dataset is clean, there are no missing, Empty values or duplicated record.

Checking Normality of Response Variable

Using this method, we obtain predictions from the model, as well as decision values from the binary classifiers.


```
library(e1071)
```

plot() - Visualizing data, support vectors and decision boundaries, if provided.

```
plot(density(df$mpg),
     main = "Milage Density Plot",
     ylab="Frequency",
     sub=paste("Skewness",round(e1071::skewness(df$mpg),2))
     )
```

```
polygon(density(df$mpg), col='#079992')
```

Milage Density Plot



Insight

- Slightly Right Skweked, which implies most of the values are posititve in nature.


#Correlation Heatmap

```
library(ggplot2)
library(reshape2)

##
## Attaching package: 'reshape2'

## The following object is masked from 'package:tidyr':
##
##   smiths

cormat <- round(cor(df),2)
melted_cormat <- melt(cormat)
ggplot(data = melted_cormat, aes(x=Var1, y=Var2, fill=value)) +
  geom_tile()
```



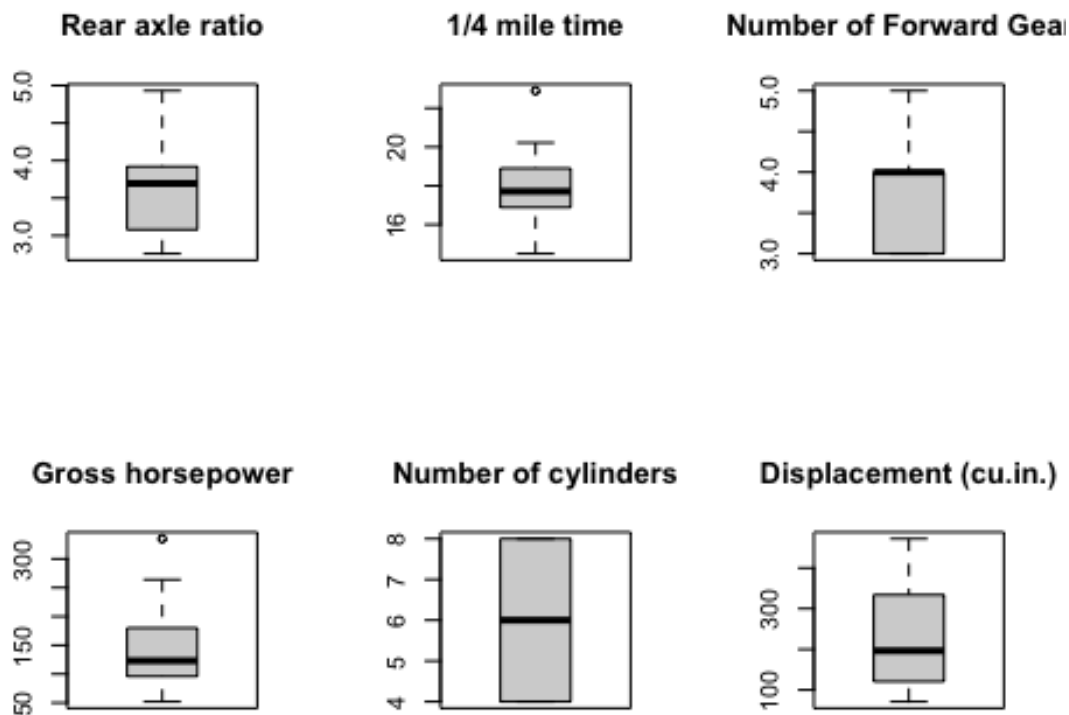
Insight

- Darker shades denotes less correlationship indications.
- Lighter shades denotes High correlationship with each variables.
- It is evident that most of the variables possess a high correlation with each other, thus we can assume multicollinearity is present.

#Checking for outliers in highly positive correlated values with Mileage

```
par(mfrow=c(2,3))

boxplot(df$drat, main = "Rear axle ratio")
boxplot(df$qsec, main = "1/4 mile time")
boxplot(df$gear, main = "Number of Forward Gears")
boxplot(df$hp, main = "Gross horsepower")
boxplot(df$cyl, main = "Number of cylinders")
boxplot(df$disp, main = "Displacement (cu.in.)")
```



Insight

- There is an outlier found in qsec(1/4 mile time) and Gross horsepower.

#Building initial model

```
X = model.matrix(mpg~. , mtcars)[, -1]
```

```
Y = mtcars$mpg
```

#Splitting the data

```
set.seed(57)
```

```
trainingRow <- sample(1:nrow(df), 0.7*nrow(df))
```

```
trainset <- df[trainingRow,]
```

```
testset <- df[-trainingRow,]
```

```
lrm <- lm(trainset$mpg~., data=trainset)
```

```
summary(lrm)
```

```
##
```

```
## Call:
```

```
## lm(formula = trainset$mpg ~ ., data = trainset)
```

```
##
```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.1759 -1.4218 -0.7548  1.0168  4.3028
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  8.55175    25.48468   0.336   0.744
## cyl         -0.07409     1.35897  -0.055   0.957
## disp         0.01402     0.02943   0.476   0.643
## hp          -0.04564     0.03718  -1.227   0.245
## drat         1.01463     2.83182   0.358   0.727
## wt          -3.66520     2.79708  -1.310   0.217
## qsec         1.22513     0.93256   1.314   0.216
## vs          -1.03224     2.89376  -0.357   0.728
## am           4.89791     2.80389   1.747   0.108
## gear        -0.76347     2.07162  -0.369   0.719
## carb         1.00937     1.36147   0.741   0.474
##
## Residual standard error: 3.006 on 11 degrees of freedom
## Multiple R-squared:  0.8924, Adjusted R-squared:  0.7945
## F-statistic:  9.12 on 10 and 11 DF,  p-value: 0.0005314
```

```
library(car)
```

```
## Loading required package: carData
```

```
##
```

```
## Attaching package: 'car'
```

```
## The following object is masked from 'package:dplyr':
```

```
##
```

```
##      recode
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
##      some
```

```
vif(lrm)
```

```
##      cyl      disp      hp      drat      wt      qsec      vs
##      am
## 14.567223 35.146018 19.377286  5.983737 20.942264  7.286081  4.928560
## 4.627183
##      gear      carb
##  4.922279 12.942380
```

Insight

- All the values are above 5, there is strong multi-collinearity present.


```

MLR_pred <- predict(lrm,testset)
compare <- cbind(actual=testset$mpg, MLR_pred)
compare

##               actual MLR_pred
## Mazda RX4      21.0 25.73138
## Mazda RX4 Wag  21.0 25.48283
## Hornet Sportabout 18.7 16.18709
## Merc 450SE     16.4 13.86304
## Merc 450SL     17.3 15.35424
## Dodge Challenger 15.5 15.86656
## Camaro Z28     13.3 12.02072
## Pontiac Firebird 19.2 15.22923
## Lotus Europa   30.4 25.48280
## Ferrari Dino   19.7 21.79969

mean (apply(compare, 1, min)/apply(compare, 1, max))

## [1] 0.8654496

RMSE = sqrt(mean((testset$mpg-MLR_pred)^2))
RMSE# calculate accuracy

## [1] 3.242592

```

Insight

- Accuracy is only 81%, which is not very efficient.

4.Choose optimum lambda value.

```

#Creating a sequence with an interval of -0.12
lambda_seq = 10^seq(3, -2, by = -.12)


# Using cross validation glmnet
ridge_model1 = cv.glmnet(X[trainingRow,], Y[trainingRow],alpha = 0,
type.measure = "mse", lambda = lambda_seq, nfolds = 5)

# Best lambda value
best_lam = ridge_model1$lambda.min
best_lam

## [1] 2.290868

plot(ridge_model1)

```



Insight

- Optimum lambda value choosed and plotted.

5.Extract the model using k-cross validation.

```
best_fit <- ridge_model1$glmnet.fit
head(best_fit)
```

```
## $a0
##      s0      s1      s2      s3      s4      s5      s6
s7
## 20.334559 20.292781 20.239225 20.171147 20.085525 19.979298 19.849725
19.694961
##      s8      s9     s10     s11     s12     s13     s14
s15
## 19.514784 19.311363 19.089707 18.858175 18.627489 18.408886 18.213989
18.048764
##      s16     s17     s18     s19     s20     s21     s22
s23
## 17.916799 17.809536 17.713541 17.605253 17.453561 17.211267 16.889578
16.443811
##      s24     s25     s26     s27     s28     s29     s30
```

```

s31
## 15.867262 15.168862 14.373356 13.507767 12.617859 11.748121 10.943290
10.235891
##      s32      s33      s34      s35      s36      s37      s38
s39
## 9.652625 9.195744 8.862197 8.633752 8.489987 8.407891 8.371597
8.363156
##      s40      s41
## 8.371917 8.390624
##
## $beta
## 10 x 42 sparse Matrix of class "dgCMatrix"

## [[ suppressing 42 column names 's0', 's1', 's2' ... ]]

##
## cyl -0.0194178748 -0.0252966251 -0.0328380896 -0.0424341586 -0.0545182791
## disp -0.0002758091 -0.0003593572 -0.0004665682 -0.0006030428 -0.0007749987
## hp -0.0004186479 -0.0005454341 -0.0007081107 -0.0009151644 -0.0011760196
## drat 0.0533808300 0.0695996484 0.0904465457 0.1170416036 0.1506476225
## wt -0.0334154670 -0.0435648435 -0.0566082128 -0.0732447062 -0.0942618997
## qsec 0.0099463820 0.0129456398 0.0167847447 0.0216558137 0.0277673440
## vs 0.0552584274 0.0719235805 0.0932558041 0.1203216130 0.1542757009
## am 0.0480936150 0.0627816340 0.0817150371 0.1059602526 0.1367494081
## gear 0.0231786989 0.0302151718 0.0392546475 0.0507776494 0.0653217181
## carb -0.0151921901 -0.0197946406 -0.0257010946 -0.0332208008 -0.0426979387
##
## cyl -0.0695388873 -0.087907973 -0.109928221 -0.135698629 -0.165015766
## disp -0.0009888889 -0.001250704 -0.001564942 -0.001933281 -0.002353179
## hp -0.0015004572 -0.001897588 -0.002374333 -0.002933503 -0.003571833
## drat 0.1926063151 0.244220147 0.306569016 0.380269080 0.465207714
## wt -0.1204958347 -0.152757675 -0.191720764 -0.237773900 -0.290863548
## qsec 0.0353274579 0.044516063 0.055444978 0.068109946 0.082345031
## vs 0.1962638343 0.247258973 0.307821869 0.377800951 0.456021858
## am 0.1754431571 0.223451530 0.282103430 0.352467645 0.435148965
## gear 0.0834498460 0.105691965 0.132453336 0.163890040 0.199763730
## carb -0.0544910930 -0.068936941 -0.086296270 -0.106686059 -0.130010072
##
## cyl -0.197283638 -0.231536306 -0.266464233 -0.300532751 -0.332346390
## disp -0.002816699 -0.003310496 -0.003816427 -0.004313505 -0.004781762
## hp -0.004278610 -0.005035651 -0.005819007 -0.006602703 -0.007362891
## drat 0.560329707 0.663541512 0.771838464 0.881686019 0.989424926
## wt -0.350375086 -0.415083181 -0.483257506 -0.552919185 -0.622047734
## qsec 0.097799533 0.113949489 0.130172146 0.145876393 0.160608946
## vs 0.540065290 0.626210735 0.709690715 0.785227191 0.847558044
## am 0.530113606 0.636590523 0.753147026 0.877942427 1.008998300
## gear 0.239303462 0.281113209 0.323170856 0.362934380 0.397554449
## carb -0.155910631 -0.183762893 -0.212736212 -0.241908068 -0.270379880
##
## cyl -0.360528450 -0.384297972 -0.402843612 -0.415951073 -0.423452517

```

```

## disp -0.005203051 -0.005565010 -0.005860181 -0.006085898 -0.006242758
## hp -0.008082416 -0.008750542 -0.009369868 -0.009946048 -0.010492122
## drat 1.091995922 1.186915651 1.273366396 1.351086594 1.420900908
## wt -0.689079862 -0.752763379 -0.812840558 -0.869349465 -0.923011403
## qsec 0.174252289 0.186977416 0.199466299 0.212655588 0.227820785
## vs 0.892306277 0.916099663 0.917099338 0.894707619 0.849574240
## am 1.144739301 1.284064223 1.426749223 1.573181967 1.724338051
## gear 0.424144323 0.440141829 0.443347589 0.432389871 0.406596869
## carb -0.297386038 -0.322313670 -0.344555506 -0.363509808 -0.378370411
##
## cyl -0.425167325 -0.42005450 -0.410120588 -0.393896701 -0.371281213
## disp -0.006333105 -0.00635348 -0.006332249 -0.006253825 -0.006116423
## hp -0.011026914 -0.01158088 -0.012176779 -0.012846804 -0.013629891
## drat 1.484296094 1.54371974 1.598688328 1.650251164 1.697789077
## wt -0.975109142 -1.02826126 -1.081801192 -1.139297646 -1.202947602
## qsec 0.246460918 0.27054145 0.300209136 0.337137522 0.381754764
## vs 0.783255381 0.69856011 0.594295394 0.475105213 0.343307682
## am 1.881503356 2.04674980 2.218420823 2.398691570 2.587386755
## gear 0.365963926 0.31086667 0.242390790 0.162379925 0.072948414
## carb -0.388033773 -0.39096250 -0.385571567 -0.370499385 -0.344427489
##
## cyl -0.342463622 -0.30873866 -0.270348779 -0.229806214 -0.188672924
## disp -0.005911888 -0.00561437 -0.005202658 -0.004637032 -0.003898958
## hp -0.014564976 -0.01566547 -0.016983067 -0.018515083 -0.020283937
## drat 1.739537496 1.77200174 1.793661659 1.801191849 1.793749442
## wt -1.274860253 -1.35695787 -1.451703556 -1.560721178 -1.684862780
## qsec 0.433804547 0.49221963 0.555823870 0.622550805 0.690414395
## vs 0.201754515 0.05567978 -0.091622500 -0.234074095 -0.368759878
## am 2.783544366 2.98484806 3.188974795 3.391538649 3.588815085
## gear -0.023108235 -0.12124949 -0.218706221 -0.310809428 -0.395541062
## carb -0.306557135 -0.25743708 -0.196781736 -0.126129760 -0.046146270
##
## cyl -0.14959444 -0.114321523 -0.0859363880 -0.0640708118 -0.049428577
## disp -0.00296822 -0.001852804 -0.0005431506 0.0008965152 0.002426298
## hp -0.02226644 -0.024430897 -0.0266952816 -0.0290085396 -0.031286481
## drat 1.77054264 1.732788652 1.6811238095 1.6198149498 1.551854476
## wt -1.82401345 -1.976241190 -2.1402540939 -2.3104968329 -2.482769028
## qsec 0.75729715 0.821420677 0.8812996172 0.9360143581 0.984986034
## vs -0.49177145 -0.601396374 -0.6940612888 -0.7717114703 -0.834179069
## am 3.77642780 3.951100196 4.1088358846 4.2489768787 4.370390310
## gear -0.47057546 -0.535317980 -0.5882799241 -0.6313053773 -0.664915963
## carb 0.04134609 0.134470039 0.2303378572 0.3267476137 0.420873237
##
## cyl -0.040722148 -0.036836911 -0.036072576 -0.037993027 -0.040630461
## disp 0.003967958 0.005458692 0.006830848 0.008074849 0.009127739
## hp -0.033461559 -0.035471692 -0.037271014 -0.038840130 -0.040163313
## drat 1.481742944 1.413049417 1.349248180 1.291229019 1.241674091
## wt -2.650541028 -2.808574581 -2.951544054 -3.078901358 -3.186057739
## qsec 1.027968753 1.064995577 1.096234180 1.122274794 1.143304513
## vs -0.883811415 -0.922412693 -0.952499479 -0.974827201 -0.992625860

```

```
## am      4.473871578  4.560469250  4.631924264  4.689726505  4.736066941
## gear   -0.690910987 -0.710635426 -0.725678449 -0.736618641 -0.745245241
## carb    0.510141494  0.592324213  0.665634193  0.729630910  0.783468884
##
## cyl    -0.04406097 -0.04759208
## disp    0.01002471  0.01075943
## hp      -0.04126095 -0.04214877
## drat     1.19944143  1.16467932
## wt      -3.27637387 -3.34992102
## qsec     1.16026691  1.17361996
## vs      -1.00570722 -1.01567187
## am       4.77257000  4.80094436
## gear    -0.75148098 -0.75619184
## carb     0.82821911  0.86447869
##
## $df
## [1] 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10
10 10
## [26] 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10 10
##
## $dim
## [1] 10 42
##
## $lambda
## [1] 1.000000e+03 7.585776e+02 5.754399e+02 4.365158e+02 3.311311e+02
## [6] 2.511886e+02 1.905461e+02 1.445440e+02 1.096478e+02 8.317638e+01
## [11] 6.309573e+01 4.786301e+01 3.630781e+01 2.754229e+01 2.089296e+01
## [16] 1.584893e+01 1.202264e+01 9.120108e+00 6.918310e+00 5.248075e+00
## [21] 3.981072e+00 3.019952e+00 2.290868e+00 1.737801e+00 1.318257e+00
## [26] 1.000000e+00 7.585776e-01 5.754399e-01 4.365158e-01 3.311311e-01
## [31] 2.511886e-01 1.905461e-01 1.445440e-01 1.096478e-01 8.317638e-02
## [36] 6.309573e-02 4.786301e-02 3.630781e-02 2.754229e-02 2.089296e-02
## [41] 1.584893e-02 1.202264e-02
##
## $dev.ratio
## [1] 0.06108987 0.07915888 0.10204834 0.13070259 0.16603686 0.20879069
## [7] 0.25931876 0.31734285 0.38172109 0.45033252 0.52018246 0.58780504
## [13] 0.64985860 0.70377117 0.74818554 0.78300796 0.80915562 0.82815348
## [19] 0.84165396 0.85118415 0.85799058 0.86302353 0.86690343 0.87007810
## [25] 0.87282648 0.87530903 0.87758983 0.87972123 0.88169272 0.88350751
## [31] 0.88515014 0.88661178 0.88788129 0.88895987 0.88984982 0.89055920
## [37] 0.89110407 0.89150530 0.89179271 0.89198805 0.89211952 0.89220514
```

6. Build the final model and interpret.

```
linRidgeMod = linearRidge(trainset$mpg ~ ., data = trainset)
predicted = predict(linRidgeMod, testset) # predict on test data
compare1 = cbind (actual=testset$mpg, predicted)

mean (apply(compare1, 1, min)/apply(compare1, 1, max))
```

```
## [1] 0.9029484

summary(linRidgeMod)

##
## Call:
## linearRidge(formula = trainset$mpg ~ ., data = trainset)
##
##
## Coefficients:
##              Estimate Scaled estimate Std. Error (scaled) t value (scaled)
## (Intercept) 14.742140              NA              NA              NA
## cyl         -0.326153        -2.753488          3.321952          0.829
## disp        -0.005767        -3.492460          2.841631          1.229
## hp          -0.015200        -5.409057          3.161192          1.711
## drat         1.757967         4.564728          3.400698          1.342
## wt          -1.322165        -6.502449          3.091123          2.104
## qsec         0.468642         4.077506          3.255874          1.252
## vs           0.103222         0.238042          3.398804          0.070
## am           2.909703         6.710119          3.236733          2.073
## gear        -0.086758        -0.279297          3.248383          0.086
## carb        -0.277200        -2.201791          3.031816          0.726
##              Pr(>|t|)
## (Intercept)              NA
## cyl              0.4072
## disp              0.2191
## hp              0.0871 .
## drat              0.1795
## wt              0.0354 *
## qsec              0.2104
## vs              0.9442
## am              0.0382 *
## gear              0.9315
## carb              0.4677
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Ridge parameter: 0.1286921, chosen automatically, computed using 3 PCs
##
## Degrees of freedom: model 5.755 , variance 4.066 , residual 7.443
```

Creating another model with only significant values.

```
linRidgeMod = linearRidge(trainset$mpg ~ ., data = trainset[, c(6,10,11)])
predicted1 = predict(linRidgeMod, testset) # predict on test data
compare2 = cbind (actual=testset$mpg, predicted1)

mean (apply(compare2, 1, min)/apply(compare2, 1, max))

## [1] 0.9464945
```

```
summary(linRidgeMod)

##
## Call:
## linearRidge(formula = trainset$mpg ~ ., data = trainset[, c(6,
##      10, 11)])
##
##
## Coefficients:
##      Estimate Scaled estimate Std. Error (scaled) t value (scaled)
## (Intercept)  27.642              NA              NA              NA
## wt          -3.316          -16.308           4.659           3.500
## gear         1.924           6.193           4.089           1.514
## carb        -1.395          -11.082           4.262           2.600
##      Pr(>|t|)
## (Intercept)      NA
## wt          0.000465 ***
## gear         0.129929
## carb        0.009319 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Ridge parameter: 0.04208334, chosen automatically, computed using 2 PCs
##
## Degrees of freedom: model 2.698 , variance 2.457 , residual 2.939

RMSE = sqrt(mean((testset$mpg-predicted1)^2))
RMSE

## [1] 1.389815
```

Insight

- The accuracy has increased from 75% to 88%.
- Root-Mean-Square Error has decreased from 3.242 to 1.389